

Value Iteration

Joshua Tsang

August 26, 2024

Contents

1	Foundational Concepts	1
1.1	Markov Reward Process vs Markov Decision Process	3
2	Playing with Episodes of an MDP	4
3	State Value Function and Bellman's Equation	5
3.1	Evaluating the state value function for an example MDP . . .	6
4	The Action Value Function, Q-Value Tables and Optimal Policies	8
5	Value Iteration	8
6	Policy Iteration	8

1 Foundational Concepts

It is instructive initially discuss the foundational concepts in Markov Decision Problems (MDP) and Reinforcement Learning (RL). Consider the problem shown in Figure 1 where an agent is located in one of the grid positions.




Figure 1: 1D Grid World MDP with an assigned reward of $+10$ in state s_{10} , termed the terminal state. The agent currently sits in state s_4 and the goal of the problem is for the agent to take actions to maximise its reward. At each state other than the terminal state, there is a negative reward of -1 . Drawn with draw.io and saved in Google Drive as 1D_grid_world.drawio.

- **State Space and State:** The finite State Space of a system is denoted S where a state s denotes a certain configuration of the system i.e. $s \in S$. For example, in Figure 1 the agent is currently in state s_4 , with the neighbouring states, $\{s'\}$, being s_3 and s_5 .
- **Actions:** For each state, s_i , that the system resides in a set of available actions, $\{a_k\}$, can be taken that transition the system to another state, s_j . For Figure 1 the available actions at state s_4 are **(move left)** and **(move right)**. One could think of a graph where nodes are states and the edges are actions.
- **Policy:** A policy, denoted π , is a distribution over actions given states. It is often written as a function $\pi(a|s)$ or $\pi(s, a)$ where it is formally defined as:

$$\pi(a|s) = \pi(s, a) = P(A_t = a | S_t = s) \quad (1)$$

which is the probability of taking action a given the system is in state s at time step t . As such, $\pi(s, a)$ returns a probability value in the range $[0, 1]$. The goal of RL is to learn the optimal policy $\pi(s, a)$ to maximise future rewards/returns (explained next). As a foreshadow, *optimal* policy functions are often expressed as an argmax_a as follows:

$$\begin{aligned} \pi_*(s, a) &= \operatorname{argmax}_a \sum_{s'} \sum_r P(s', r | s, a) [r + \gamma v_\pi(s')] \\ &= \operatorname{argmax}_a q_*(s, a) \end{aligned} \quad (2)$$

i.e. the optimal policy function returns the action a that transitions the agent to the neighbouring state s' that yields the maximum transition-probability-weighted sum of the returns, or in other words, the maximum *action value* as discussed around equation (14).

- **Reward:** Numerical reward can be associated with certain states and actions, represented as R_t for the current time step. In the case of Figure 1, there is a fixed reward of +10 at state s_{10} and the other states have a reward of -1. The goal of RL can be considered to be finding the sequence of actions that maximise the sum of rewards to achieve a successful episode i.e. the agent reaches the terminal state s_{10} .
- **Episodes:** MDPs have the characteristic of being finite i.e. they end within a finite number of time steps, denoted T . More specifically, an episode is a sequence of states, actions and rewards that start at some state and end in the terminal state.
- **Return or "Future Gains":** Usually denoted G_t it is the sum of rewards gained by the agent *after* time step t , define as:

$$G_t = R_{t+1} + R_{t+2} + \dots + R_T \quad (3)$$

Before going further, it's worth appreciating how important this concept is to the definition of state value functions, $v(s)$, which try to capture the expected rewards to be gained by being in that state, s .

It is possible to introduce a discount rate $\gamma \in [0, 1]$ that diminishes the rewards gained further in the future:

$$\begin{aligned} G_t &= R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots + \gamma^{T-1-t} R_T \\ &= \sum_{k=t+1}^T \gamma^{k-1-t} R_k \end{aligned} \quad (4)$$

note how the power of γ in the final term is $T-1-t$ and *not* T , a little thought makes this clear (for example, try write down G_4 where $T = 8$). Finally, we note that the discounted expression can be rewritten into a recursive expression somewhat similar to Bellman's equation:

$$\begin{aligned} G_t &= R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots + \gamma^{T-1-t} R_T \\ &= R_{t+1} + \gamma(R_{t+2} + \gamma R_{t+3} + \dots + \gamma^{T-t} R_T) \\ &= R_{t+1} + \gamma G_{t+1} \end{aligned} \quad (5)$$

Expression (5) links the current return, G_t , to the return at the next time step, G_{t+1} .

1.1 Markov Reward Process vs Markov Decision Process

It is instructive to discuss Markov Reward Process (MRP) which are simpler than a Markov Decision Process (MDP). This discussion helps consolidate concepts surrounding transition probabilities, policies and action values. In short, an MDP is essentially an MRP *with actions* and thus an MRP will still have states, rewards and transition probabilities. The transition probabilities encapsulate the dynamics of the process or the effects of the environment.

Consider the MRP shown in Figure 2 which describes a 'slippery world' MRP where the transition probabilities of moving to adjacent states is only successful with probability 0.8, with a 0.2 chance of remaining in the same state.

The state transition matrix, \mathbf{P} , for the MRP can be written:

$$\begin{aligned} \mathbf{P} &= \begin{pmatrix} P(s_1|s_1) & P(s_2|s_1) & P(s_3|s_1) & \dots & P(s_N|s_1) \\ P(s_1|s_2) & P(s_2|s_2) & P(s_3|s_2) & \dots & P(s_N|s_2) \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ P(s_1|s_N) & P(s_2|s_N) & P(s_3|s_N) & \dots & P(s_N|s_N) \end{pmatrix} \\ &= \begin{pmatrix} 0.2 & 0.8 & 0.0 & 0.0 & 0.0 \\ 0.8 & 0.2 & 0.8 & 0.0 & 0.0 \\ 0.0 & 0.8 & 0.2 & 0.8 & 0.0 \\ 0.0 & 0.0 & 0.8 & 0.2 & 0.8 \\ 0.0 & 0.0 & 0.0 & 0.8 & 0.2 \end{pmatrix} \end{aligned} \quad (6)$$




Figure 2: The Slippery World MRP where there's a 0.2 chance of remaining in the present state as the agent can slip and fall, making no movement progress. There are still rewards associated with the states, hence the name Markov *Reward* Process. Drawn with draw.io and saved in Google Drive.

Knowing these state transition probabilities allows trajectories to be sampled and expect returns to be calculated. Like an MDP, state values for each state, $v(s)$, can be computed. In fact, the state values can be computed analytically using the linear equation:

$$\begin{pmatrix} v(s_1) \\ v(s_2) \\ \vdots \\ v(s_5) \end{pmatrix} = \begin{pmatrix} r(s_1) \\ r(s_2) \\ \vdots \\ r(s_5) \end{pmatrix} + \gamma \begin{pmatrix} P(s_1|s_1) & P(s_2|s_1) & P(s_3|s_1) & \dots & P(s_5|s_1) \\ P(s_1|s_2) & P(s_2|s_2) & P(s_3|s_2) & \dots & P(s_5|s_2) \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ P(s_1|s_5) & P(s_2|s_5) & P(s_3|s_5) & \dots & P(s_5|s_5) \end{pmatrix} \begin{pmatrix} v(s_1) \\ v(s_2) \\ \vdots \\ v(s_5) \end{pmatrix}$$

$$\mathbf{v} = \mathbf{r} + \gamma \mathbf{P}\mathbf{v}$$
(7)

which allows the state values to be solved using some linear algebra as:

$$\mathbf{v} = (\mathbf{I} - \gamma \mathbf{P})^{-1} \mathbf{r}$$
(8)

Solving the matrix is roughly $O(N^3)$ where N is the number of states.

2 Playing with Episodes of an MDP

Recall that episodes are a sequence of states, actions and rewards that end in the terminal state. Each episode yields a return/gain at T , which is G_T . To clarify concepts and develop some intuitive, consider a few episodes for the Grid World system in Figure 1. The actions are $a \in [\text{LEFT}, \text{RIGHT}]$

- Episode Example 1:

$$\begin{aligned} [S_{t=1} = s_4, A_{t=1} = \text{RIGHT}, G_{t=1} = -1] \\ [S_{t=2} = s_5, A_{t=2} = \text{RIGHT}, G_{t=2} = -2] \\ [S_{t=3} = s_6, A_{t=3} = \text{RIGHT}, G_{t=3} = -3] \\ [S_{t=4} = s_7, A_{t=4} = \text{RIGHT}, G_{t=4} = -4] \end{aligned}$$

$$\begin{aligned}
&[S_{t=5} = s_8, A_{t=5} = \text{RIGHT}, G_{t=5} = -5] \\
&[S_{t=6} = s_9, A_{t=6} = \text{RIGHT}, G_{t=6} = -6] \\
&[S_{t=7} = s_10, A_{t=7} = \text{RIGHT}, G_{t=7} = +4]
\end{aligned}$$

Note that this episode yields $G_T = +4$.

- Episode Example 2:

$$\begin{aligned}
&[S_{t=1} = s_4, A_{t=1} = \text{RIGHT}, G_{t=1} = -1] \\
&[S_{t=2} = s_5, A_{t=2} = \text{RIGHT}, G_{t=2} = -2] \\
&[S_{t=3} = s_6, A_{t=3} = \text{RIGHT}, G_{t=3} = -3] \\
&[S_{t=4} = s_7, A_{t=4} = \text{RIGHT}, G_{t=4} = -4] \\
&[S_{t=5} = s_8, A_{t=5} = \text{LEFT}, G_{t=5} = -5] \\
&[S_{t=6} = s_7, A_{t=6} = \text{RIGHT}, G_{t=6} = -6] \\
&[S_{t=7} = s_8, A_{t=7} = \text{RIGHT}, G_{t=7} = -7] \\
&[S_{t=8} = s_9, A_{t=8} = \text{RIGHT}, G_{t=8} = -8] \\
&[S_{t=9} = s_10, A_{t=9} = \text{RIGHT}, G_{t=9} = +2]
\end{aligned}$$

Note that this episode yields $G_T = +2$, it is lower than Episode Example 1 because at $t = 5$ it backtracks by going LEFT which caused it to require a longer path, and thus suffer more from the $R_t = -1$ penalties for the non-terminal states.

3 State Value Function and Bellman's Equation

An important concept in RL is the state value function, $v_\pi(s)$, which captures the expected return from being in state s . It is critical to realise that the subscript π in $v_\pi(s)$ means it must be evaluated for a given policy $\pi(a|s)$ since $v_\pi(s)$ is an expectation value (read: average) of potential returns by being in state s , which in turn depends on the expected return from the neighbouring states, $v_\pi(s')$. In fact, this is why the Bellman equation in equation (11) has $v_\pi(s)$ on the LHS and $v_\pi(s')$'s on the RHS.

Since the policy $\pi(a|s)$ encapsulates the current policy's probability of taking action a given the current state is s (recall equation 1), it thus influences the expected returns of being in the current state s . In fact, it's prudent to think of the state value function as a function that ultimately links the state values of all states together.

The state value function, $v_\pi(s)$ of an MDP is the expected return starting from state s and following the policy π :

$$\begin{aligned}
v_\pi(s) &= \mathbb{E}[G_t|s] \\
&= \sum_a \pi(a|s) \mathbb{E}[G_t|s, a]
\end{aligned} \tag{9}$$

This is fairly intuitive, the state value is equal to the sum of the expected returns from all available actions a at state s weighted by the present policy probability distribution, π .

Using Equation (5) we can rewrite $\mathbb{E}[G_t|s, a]$ as:

$$\begin{aligned}\mathbb{E}[G_t|s, a] &= \mathbb{E}[R_{t+1} + \gamma G_{t+1}|s, a] \\ &= \mathbb{E}[R_{t+1} + \gamma v_\pi(S_{t+1})|s, a] \\ &= \sum_{s'} \sum_r P(s', r|s, a)[r + \gamma v_\pi(s')]\end{aligned}\tag{10}$$

This is in fact the action value function as discussed in the next section. The final expression after substitution back into equation (9) is:

$$v_\pi(s) = \sum_a \pi(a|s) \sum_{s'} \sum_r P(s', r|s, a)[r + \gamma v_\pi(s')]\tag{11}$$

Equation (11) allows us to calculate the expected return value for each state in the MDP given the current policy, π , which may not yet be optimal. We haven't yet discussed optimal policies at all and equation (11) simply calculates state values for a given policy $\pi(a|s)$. It is the goal of RL to find the optimal policy, π_* , via methods like General Policy Iteration (GPI), Policy Iteration or Value Iteration to discover optimal policies.

Note that the probabilities $P(s', r|s, a)$ are not to be confused with the policy $\pi(s, a)$ but are transition probabilities intrinsic to the MDP that capture the dynamics of the environment. Note that for an MRP the state transition probabilities are merely $P(s'|s)$ while for an MDP (which includes actions) it is $P(s', r|s, a)$ which encapsulates probabilities to transition into s' for a given action a . For instance, we could introduce a strong fan effect that blows left in state s_4 of the MDP in Figure 1 where there's now a 0.2 chance of the agent losing its footing/grip and will be blown back to state s_3 . These environmental dynamics/effects are captured in the transition probabilities of the MDP. So in this case, if we're in state s_4 then $P(s_5, r|s_4, a = \text{RIGHT}) = 0.8$ and $P(s_3, r|s_4, a = \text{RIGHT}) = 0.2$ i.e. there's a 0.2 chance of transitioning into the state in the opposite direction. The goal of RL is to deduce the optimal policy taking into account these environmental/dynamic effects.

Now is a good time to mention model-based and model-free processes. A model-based process has full knowledge of the transition probabilities $P(s', r|s, a)$, which although introduce a probabilistic element to the process, these transitions are known and the MDP is well-defined.

(Talk about Gambler's problem and the transition probabilities in that.)

3.1 Evaluating the state value function for an example MDP

Consider the 'Slippery World' MRP shown in Figure 2 and consider that we promote it to an MDP complete with actions. At each state, there are 2

actions $a \in \{a_1, a_2\}$ where a_1 and a_2 represent the actions move LEFT and move RIGHT respectively. Since the MDP has $S = 5$ states and each state has $A = 2$ actions, the total number of possible policies is $A^S = 2^5 = 32$. Each state only has one ‘good’ action so one can imagine each policy being like a binary sequence: 01110 for example means $[a_1, a_2, a_2, a_2, a_1]$ and it’s easy to imagine there are 32 possible combinations of these actions.

Let’s try and evaluate the state value function equation (11) for state s_4 given the policy defined as follows:

$$\begin{bmatrix} (\pi(a_1|s_1), \pi(a_2|s_1)) \\ (\pi(a_1|s_2), \pi(a_2|s_2)) \\ (\pi(a_1|s_3), \pi(a_2|s_3)) \\ (\pi(a_1|s_4), \pi(a_2|s_4)) \\ (\pi(a_1|s_5), \pi(a_2|s_5)) \end{bmatrix} = \begin{bmatrix} (0.4, 0.6) \\ (0.4, 0.6) \\ (0.4, 0.6) \\ (0.4, 0.6) \\ (0.4, 0.6) \end{bmatrix} \quad (12)$$

So the policy here is a 0.4 and 0.6 chance of taking action to move LEFT and RIGHT respectively, it effectively tries to do a 40/60 random walk in this slippery world but note it might not be able to actually do a perfect random walk due to the slippery dynamics encoded into the transition probabilities $P(s', r|s, a)$. Computing $v_\pi(s_4)$ with a discount factor of $\gamma = 0.5$:

$$\begin{aligned} v_\pi(s_4) &= \pi(a_1|s_4)[P(s_4, r|s_4, a_1)(r + \gamma v_\pi(s_4)) + P(s_3, r|s_4, a_1)(r + \gamma v_\pi(s_3))] \\ &\quad + \pi(a_2|s_4)[P(s_4, r|s_4, a_2)(r + \gamma v_\pi(s_4)) + P(s_5, r|s_4, a_2)(r + \gamma v_\pi(s_5))] \\ &= 0.4[0.2(-1 + 0.5 \times -1) + 0.8(-1 + 0.5 \times -1)] \\ &\quad + 0.6[0.2(-1 + 0.5 \times -1) + 0.8(+10 + 0.5 \times +10)] \\ &= -0.36 + 7.38 \\ &= 7.02 \end{aligned} \quad (13)$$

A similar computation can be performed for all states in the MDP. Note that a different policy π would give a different value for $v_\pi(s_4)$. For example, if $\pi(a_1|s_4) = 1.0$ and $\pi(a_2|s_4) = 0.0$ then none of s_5 ’s significant +10 reward will be included in the value of $v_\pi(s_4)$ yielding a much lower state value. It is clear that for each state, there exists a policy that would maximise their state values. For s_4 it is clear to let $\pi(a_1|s_4) = 0.0$ and $\pi(a_2|s_4) = 1.0$ so the agent always tries to move RIGHT to claim that significant +10 reward.

Note what might happen through repeated applications/sweeps of equation (11) to all the states of the MDP. It intuitive to see that the fixed +10 reward in state s_5 will gradually ‘trickle’ down into the state values of the states on the left.

4 The Action Value Function, Q-Value Tables and Optimal Policies

The action value function, denoted q_π , is defined as the inner sums in Equation (11):

$$q_\pi(s, a) = \sum_{s'} \sum_r P(s', r | s, a) [r + \gamma v_\pi(s')] \quad (14)$$

where s' are the neighbouring states to the current state, s . Do not confuse state values functions and action value functions. Action value functions include the action in the argument, naturally! It might be helpful to imagine equation (14) computed in matrix for a particular MDP in the following manner:

$$\begin{bmatrix} q(s_1, a_1) & q(s_1, a_2) & \dots & q(s_1, a_K) \\ q(s_2, a_1) & q(s_2, a_2) & \dots & q(s_2, a_K) \\ \vdots & \vdots & \ddots & \vdots \\ q(s_N, a_1) & q(s_N, a_2) & \dots & q(s_N, a_K) \end{bmatrix} \quad (15)$$

where N is the number of states and K is the number of actions per state. Matrix (15) is often referred to as a ‘Q-Value Table’.

In Q-Learning, once the Q-Value Table has been constructed, it can be used by the agent to determine the next action for a given state by selecting the action that yields the maximum q value. So say for s_4 in the Slippery World MDP, the agent will look at the row for s_4 pick the action a yielding the greatest $q(s_4, a)$. I am wondering how matrix (15) can be computed without first knowing the optimal policy, and I think the answer is that Q-Learning is a form of off-policy RL which I have not yet covered.

Equation (14) allows the optimal policy, π_* , to be written succinctly as:

$$\pi_*(s) = \operatorname{argmax}_a q_\pi(s, a) \quad (16)$$

You will notice that ‘optimal’ expressions usually have some form of \max_s or argmax_a in it.

5 Value Iteration

Value iteration allows one to assign state values to each state through repeated, iterative application of the value function, and the optimal policy is extracted at the end.

6 Policy Iteration


```

Iteration = 0, V_new = {1: -1.0, 2: -1.0, 3: -1.0, 4: -1.0, 5: -1.0, 6: -1.0, 7: -1.0, 8: -1.0, 9: -1.0, 10: 1.0}
Iteration = 1, V_new = {1: -2.0, 2: -2.0, 3: -2.0, 4: -2.0, 5: -2.0, 6: -2.0, 7: -2.0, 8: -2.0, 9: 0.0, 10: 1.0}
Iteration = 2, V_new = {1: -3.0, 2: -3.0, 3: -3.0, 4: -3.0, 5: -3.0, 6: -3.0, 7: -3.0, 8: -1.0, 9: 0.0, 10: 1.0}
Iteration = 3, V_new = {1: -4.0, 2: -4.0, 3: -4.0, 4: -4.0, 5: -4.0, 6: -4.0, 7: -2.0, 8: -1.0, 9: 0.0, 10: 1.0}
Iteration = 4, V_new = {1: -5.0, 2: -5.0, 3: -5.0, 4: -5.0, 5: -5.0, 6: -3.0, 7: -2.0, 8: -1.0, 9: 0.0, 10: 1.0}
Iteration = 5, V_new = {1: -6.0, 2: -6.0, 3: -6.0, 4: -6.0, 5: -4.0, 6: -3.0, 7: -2.0, 8: -1.0, 9: 0.0, 10: 1.0}
Iteration = 6, V_new = {1: -7.0, 2: -7.0, 3: -7.0, 4: -5.0, 5: -4.0, 6: -3.0, 7: -2.0, 8: -1.0, 9: 0.0, 10: 1.0}
Iteration = 7, V_new = {1: -8.0, 2: -8.0, 3: -6.0, 4: -5.0, 5: -4.0, 6: -3.0, 7: -2.0, 8: -1.0, 9: 0.0, 10: 1.0}
Iteration = 8, V_new = {1: -9.0, 2: -7.0, 3: -6.0, 4: -5.0, 5: -4.0, 6: -3.0, 7: -2.0, 8: -1.0, 9: 0.0, 10: 1.0}
Iteration = 9, V_new = {1: -8.0, 2: -7.0, 3: -6.0, 4: -5.0, 5: -4.0, 6: -3.0, 7: -2.0, 8: -1.0, 9: 0.0, 10: 1.0}
Iteration = 10, V_new = {1: -8.0, 2: -7.0, 3: -6.0, 4: -5.0, 5: -4.0, 6: -3.0, 7: -2.0, 8: -1.0, 9: 0.0, 10: 1.0}
Iteration = 11, V_new = {1: -8.0, 2: -7.0, 3: -6.0, 4: -5.0, 5: -4.0, 6: -3.0, 7: -2.0, 8: -1.0, 9: 0.0, 10: 1.0}

```

Figure 3: Iterations of the value iteration algorithm.